

RISHIN S PRADEEP

Senior Data Architect | Snowflake & dbt Technical Authority | ELT Architecture | Data Strategy

■ +91 8825959211 ■ talktorishin94@gmail.com ■ linkedin.com/in/rishin-s-pradeep-28420ba7

■ Bengaluru, INDIA



SUMMARY

Results-driven **Senior Data Architect** with 8+ years of end-to-end experience architecting enterprise-scale ELT platforms, **Snowflake** data ecosystems, and **dbt** modeling frameworks. Proven **technical lead** of large global engineering teams (70+ members), delivering zero-downtime migrations, **CI/CD governance**, and production-grade decision products. Deep expertise in **Snowflake Virtual Warehouse tuning**, **RBAC frameworks**, **query performance optimization**, **cost engineering**, and **semantic layer design**. Hands-on **dbt** authority across modularity, testing, documentation, macros, exposures, and governance. Skilled at bridging technical depth with business clarity—mentoring L1/L2 engineers, leading executive architecture reviews, and shaping long-term data platform roadmaps.

EXPERIENCE

Senior Data Architect

Quantiphi Analytics

■ 08/2025 - Present ■ India

Technical lead of a 70-member global data engineering and DevOps coalition, executing enterprise-scale **Snowflake** cloud modernization and **Data Vault 2.0** architecture.

- **Led and mentored** a 70-member global data engineering team, conducting **design reviews**, **code reviews**, and **architectural coaching** to drive engineering maturity—executing a flawless, zero-downtime deprecation of a monolithic on-premise **Exasol** ecosystem (billions of CRM records from Salesforce & Siebel) into **Snowflake**.
- Architected dynamic Python orchestrators triggering **Ephemeral GitHub Actions Runners** via REST API webhooks; runners queried Exasol locally, streamed data into **.csv.gz buffers**, and pushed asynchronously to **Snowflake External Stages** (cloud storage buckets)—entirely bypassing local memory bottlenecks at scale.
- Engineered an automated, idempotent **CI/CD pipeline** with dynamic **Snowflake schema mirroring** using **Zero-Copy Cloning** for DDL instantiation, regex-driven column sanitization, cross-database type casting, and multi-threaded exponential-backoff retry logic to overcome firewall idle-timeouts during heavy **COPY INTO** extraction cycles.
- Designed the full downstream architecture using **Airflow (Cloud Composer)** as master scheduler, triggering containerized **dbt** nodes to transform External Stage dumps into a **Data Vault 2.0** ecosystem (Bronze/Silver/Gold tiers), leveraging **dbt modularity**, **dbt testing**, **dbt documentation**, and **dbt exposures** to natively optimize **Snowflake micro-partitioning**.
- Engineered **Snowflake Virtual Warehouse** scale-up/scale-down orchestration to parallel-ingest 13GB+ CRM snapshots, applying **Multi-Cluster Warehouse** strategies and optimal data chunking to eliminate execution bottlenecks—**cutting global BI reporting latencies** significantly.
- **Delivered executive-level presentations** on architecture decisions, migration tradeoffs, and **CI/CD governance** strategy to senior stakeholders; coordinated **on-call schedules** and **release management** processes ensuring uninterrupted production support.

Senior Data Architect

kipi.bi

■ 01/2025 - 07/2025 ■ India

Technical authority on **Snowflake** and **dbt** for a retail analytics platform powering \$12B+ in Walmart Luminate insights.

- Served as **Snowflake technical authority**—designed **Clustering Key** strategies, **materialization policies** (table, incremental, ephemeral, view), and **Virtual Warehouse sizing** to maximize query performance and minimize compute cost for analytics-ready datasets extracted from Walmart Luminate.
- Architected and governed 350+ **modular dbt models** with full **dbt testing** (**not_null**, **unique**, **accepted_values**, **relationships**), **dbt documentation** (model descriptions, column-level docs), and **dbt exposures** to surface data lineage to business

stakeholders—powering executive KPIs.

- Reduced **dbt model run time by 60%** through **CTE refactoring**, **model materialization tuning**, **Snowflake Clustering Key** optimization, and intelligent **dbt ref() / source()** dependency graph structuring.
- Implemented **dbt macros** to standardize complex business logic across models, enforced **dbt governance** through modular staging/intermediate/mart layer separation, and engineered a **self-serve analytics framework** via dbt docs—increasing analyst productivity by 40% and reducing reliance on engineering.
- Owned the **CI/CD pipeline** for dbt using **GitHub Actions**, supporting 15+ daily deployments with **automated dbt test gates**, schema checks, and **regression testing**—reducing production incidents by 70%.
- Streamlined end-to-end pipelines using **Azure Data Factory**, achieving 99.9% SLA adherence; shortened data delivery from 3 days to under 6 hours by aligning **Snowflake** data flows with business KPI requirements.

Data Architect

Tredence Analytics

■ 12/2021 - 12/2024 ■ India

Led enterprise data integration and **ELT architecture** for ecommerce analytics, integrating Amazon SP API data into **Azure Synapse** analytical platforms.

- Architected scalable **ELT pipelines** using **Azure Synapse Serverless SQL Pool**, enabling near real-time ecommerce insights and reducing manual reporting by 90%; applied **semantic layer design** principles to align backend models with executive KPIs.
- Integrated Amazon Selling Partner (SP) API across 20+ endpoints with **hourly sync frequency**, standardized ingestion using **Parquet format** achieving 60–70% storage/compute reduction, and enabled sub-second query latency on ~500M+ record datasets.
- Designed **metadata-driven external table architecture** supporting schema evolution and rapid onboarding of new data sources; delivered 10+ executive dashboards reducing Power BI load time from 3 minutes to **<10 seconds**.
- Established **CI/CD workflows** via **Azure DevOps** for deploying metadata, schemas, and SQL logic; defined **DataOps best practices** including **observability**, **monitoring**, **alerting**, and versioning standards—reducing config errors by ~80%.
- **Mentored junior engineers** across squads, conducted **technical code reviews**, and

facilitated stakeholder sessions to translate 30+ business KPIs into optimized data models aligned with reporting goals.

Senior Data Engineer

Tredence Analytics

■ 08/2022 - 09/2023 ■ India

Designed high-performance **Azure Synapse Dedicated SQL Pool** architecture supporting finance, sales, and operations analytical workloads.

- Engineered 50+ **parameterized, metadata-driven ADF pipelines** automating 3+ TB/month ingestion; improved pipeline reusability by 70% with zero code duplication.
- Achieved **50% reduction in data load time** by optimizing **Synapse loading strategies** (PolyBase vs COPY INTO), partitioning large fact tables, and minimizing distribution skew.
- Enforced **RBAC policies**, column-level masking, and **row-level security** to ensure GDPR compliance; implemented **row-level and schema-level data quality checks** achieving ~98% reporting accuracy.
- Applied **monitoring and alerting** for ADF and Synapse jobs, improving SLA compliance from 94% to 99.8%; authored **architecture documentation** and data lineage flows reducing onboarding time by 50%.

Technology Analyst

Infosys Private Limited

■ 06/2017 - 12/2021 ■ India

Led legacy data warehouse modernization to **Azure Synapse Dedicated SQL Pool**, enabling 10x scalability and 4x faster query performance.

- Designed and deployed **Kimball-style dimensional models** (fact-dimension schemas) resulting in 30% faster report generation; translated 40+ executive KPIs into optimized data models in close partnership with Product, BI, and Finance teams.
- Built 50+ **high-performance T-SQL stored procedures** encoding complex business rules, aggregations, and hierarchies—achieving data accuracy >97%.
- Executed advanced **query performance tuning**: large table **partitioning**, **distribution strategies**, **materialized views**, and indexing—cutting query run time by 60% and infrastructure costs by 25%.
- Established **row-level security**, **RBAC policies**, and column-level masking for GDPR/PII compliance; enabled real-time self-serve analytics for 100+ users with >99.9% warehouse uptime.

LANGUAGES

English



Native

Hindi



Advanced

SNOWFLAKE EXPERTISE

- VWH Tuning & Multi-Cluster Warehouse** (Auto/Manual Scale)
- Auto-Suspend/Resume** · **Zero-Copy Cloning** (DDL & Data)
- External Stages & COPY INTO** · **Snowpipe**
- Clustering Keys** & Micro-Partition Pruning
- RBAC Frameworks** · **Row-Level Security** · Column Masking
- Dynamic Data Masking** · **Data Sharing & Secure Views**
- Query Profile & Cost Engineering**
- Materialized Views** · Result Cache · **Time Travel**
- Streams & Tasks** · **Stored Procedures & UDFs**
- Information Schema & Account Usage** Views
- Semantic Layer Design** · **Fail-Safe**

DBT EXPERTISE

- Modular Architecture**: Staging / Intermediate / Mart layers
- Materializations**: table, incremental, ephemeral, view
- Testing**: not_null, unique, accepted_values, relationships
- Custom Generic & Singular Tests** · dbt-expectations
- Macros & Jinja Templating** · **dbt Packages**
- Documentation** · Column Descriptions · **Exposures**
- Sources & Freshness Checks** · **Seeds & Snapshots**
- CI/CD**: GitHub Actions + dbt Cloud automated test gates
- ref() / source()** graph optimization · **Pre/Post Hooks**
- 350+ Production dbt Models** Delivered

SKILLS

Snowflake	dbt Core / Cloud
Apache Spark	Azure Synapse
Azure Data Factory	Databricks
Airflow / Cloud	
Composer	GitHub Actions CI/CD
Python	SQL / T-SQL
Data Vault 2.0	Terraform
Fivetran	AWS / GCP
Tableau	Power BI

EDUCATION

B.Tech/B.E. — Electrical & Electronics

Vellore Institute of Technology

■ India

12th — CBSE

10th — CISCE (ICSE/ISC)

■ India